Leakage Detection in Water-Distribution-Networks using Physics-Aware Neural Networks and Leakage Localization using a Rule-based System

Jash Shah[a,1], David Daniels[a,1], Heet Shah[a,1], Abhijeet Salunke[a,1]

[a]*Department of Computer Engineering, Sardar Patel Institute of Technology, Mumbai, India*

Abstract

Abstract example for the Results in Physics Journal. Here you provide a brief summary of the research and the results.

Keywords: keyword 1, keyword 2, keyword 3, keyword 4

1. Introduction

1.1. Motivation and Relevance

Water is one of the most essential resources for human survival, economic growth, and sustainable urban development. With rapid urbanization and increasing population density in major towns and cities, Water Distribution Networks (WDNs) have become more extensive and complex. These networks are responsible for delivering potable water from natural water bodies or reserviors to consumers, and their efficient operation is crucial to ensure water security and sustainability. The volume of leakages in WDNs varies throughout different geographies, and the primary factors affecting the amount of water lost due to leakage are the general condition of the network and the maintenance and monitoring system in place. Some places suffer from as much as 70-80% loss in water due to leakages, while where the infrastructure is secure, the lost water volume is around 7-8%. In total, around the World, the amount of non-revenue water (NRW) is around 126 billion cubic meters annually, and amounts to monetary loss of approximately 39 billion dollars(USD.)

1.2. Challenges

Monitoring systems for WDNs are primarily acoustic monitoring systems, or they rely heavily on manual inspection and are labor-intensive, time-consuming, and costly. Recent advancements in data-driven methods and sensor-based monitoring have improved detection accuracy; however, many of these approaches rely heavily on data driven and complex models that do not explicitly incorporate the underlying physical laws of fluid dynamics. As a result, their predictions may lack physical interpretability and generalization in unseen scenarios.

1.3. Contribution

Our model, which is Physics-Aware ANN (PAANN), tries to integrate physical laws into the model latent space so the predictions do not violate the equations and laws of conservation of mass and energy, and the models vastly reduce the number of false-positives which save both the monetary and human cost, predicting a leak only when the model is very sure. This is a comparatively low cost solution that can be generalized for any WDN, and our localization method predicts the leakage area using a relatively simple set of rules, not relying on expensive models for localization which are both time and compute consuming.

2. Related work

The detection and localization of leaks in Water Distribution Networks (WDNs) have been investigated through various methods. This section presents a synthesis of recent advances, organized into subsections showing the various methodologies applied, ranging from tradition data driven machine learning approaches to physics informed ones , along with the research gap motivating the present study.

2.1. Traditional Data-Driven Approaches

Classical machine learning and deep learning methods have demonstrated considerable promise in leak detection, primarily by learning patterns from historical operational data. **?** employed advanced statistical features—specifically, Autocorrelation Coefficient (Au-C) and Signal Energy (Sig-E)—extracted from accelerometer and dynamic pressure sensor measurements. Their Artificial Neural Network (ANN) achieved accuracies of 86.5% and 86.2% F1-scores for branched and looped configurations, respectively. The study highlighted the efficacy of data fusion in enhancing detection reliability; however, the approach remained purely empirical, lacking integration with underlying hydraulic principles.

Similarly, **?** investigated machine learning strategies including K-means clustering, Learning Vector Quantization (LVQ), and cluster validation techniques on data from a District Metered Area (DMA) in Stockholm, Sweden. Their methodology achieved classification rates up to 93.98%, demonstrating robustness to common sensor noise levels within ±5%. Nonetheless,

[1] All authors contributed equally to this work.

[2] {jash.shah23, david.daniels23, heet.shah223, abhijeet_salunke}.spit.ac.in

the authors acknowledged limitations in handling dynamic operational variations and the absence of physical interpretability in model predictions.

Deep learning architectures have also been explored. **?** applied Convolutional Neural Networks (CNN) and Long Short-Term Memory (LSTM) networks on the LeakDB benchmark dataset **?**, achieving leak detection accuracies ranging from 49% to 98.23% depending on network complexity. While these models captured temporal dependencies effectively, they required extensive labeled training data and exhibited reduced generalization under distribution shifts—a common challenge in purely data-driven paradigms.

2.2. Hybrid and Model-Based Methodologies

Recognizing the limitations of purely data-driven approaches, recent research has integrated hydraulic modeling with statistical or machine learning techniques. **?** proposed a hybrid framework combining EPANET hydraulic simulations with a Time-Windowed Head Bias Correction (TWHBC) mechanism and an Elastic-Net logistic regression classifier. Their approach explicitly accounted for model-data discrepancies through bias correction and achieved leak localization within predefined search radii for multiple in-field experiments. However, the methodology's reliance on accurate hydraulic model calibration and computational overhead limited its real-time applicability.

? introduced a multilayer network representation integrating topological, correlation-based, and sensor coverage graphs. By combining graph theory with statistical anomaly detection (z-score and Interquartile Range), they successfully detected leaks within 15 minutes and localized them within a 50-meter range in the BattLeDIM benchmark network. Their approach reduced the need for exhaustive hydraulic simulations by regionalizing search spaces; nevertheless, the framework struggled with incipient leaks exhibiting gradual flow increases, underscoring the necessity for continuous physical constraint enforcement.

2.3. Physics-Informed Neural Networks in WDNs

Physics-Informed Neural Networks (PINNs) represent an emerging paradigm that embeds governing partial differential equations (PDEs) directly into neural network loss functions, ensuring adherence to known physical laws. **?** pioneered the application of a 1D CNN deep autoencoder with semi-supervised learning on multivariate SCADA data from the L-TOWN benchmark network. Their PINN-based framework incorporated continuity and momentum equations via Pruned Exact Linear Time (PELT) change point detection, achieving identification of 16 out of 19 leaks and accurate localization of 13 within a 300-meter radius—without requiring a calibrated hydraulic model. The study demonstrated minimal identification lag for abrupt leaks, though incipient leaks remained challenging.

Expanding on this foundation, **?** explored PINNs for cybersecurity applications in WDNs, specifically mitigating sensor-based cyberattacks. By training surrogate models on Navier-Stokes equations, they achieved predictive accuracies exceeding 99% for velocity fields, enabling the detection of data integrity attacks within simulated environments. While conceptually innovative, the work focused on small-scale proof-of-concept scenarios and did not address scalability to operational networks.

More recently, **?** developed a digital twin framework leveraging PINNs for pipeline leakage detection in Taiwan's water infrastructure. Their model embedded continuity and momentum equations into the loss function and utilized statistical control charts (z-score and IQR) on reconstruction errors to detect anomalies. Validation on real-world datasets yielded R^2 values of 88.91% and 98.91% for small- and large-scale networks, respectively, with leak detection within 15 minutes of occurrence. Despite these successes, the framework's dependency on accurate boundary conditions and its limited handling of transient flow regimes were noted as areas requiring further investigation.

2.4. Leak-Specific Graph Learning Approaches

Recent studies have moved beyond generic physics-informed modeling toward frameworks explicitly tailored for leak detection and localization in Water Distribution Networks (WDNs), integrating graph-based learning with hydraulic constraints to better capture leak-induced anomalies.

? proposed a physics-informed graph neural network framework that embeds mass conservation and pressure–flow relationships directly into the learning process for leak localization. By representing the WDN as a graph and propagating leak-induced pressure residuals across adjacent nodes, the model exploited spatial correlations inherent in the network topology to improve localization accuracy. The incorporation of physics-based loss terms significantly reduced false positives compared to purely data-driven GNN baselines, particularly under moderate sensor noise conditions. Experimental validation demonstrated improved robustness to measurement uncertainty and partial observability. However, the approach primarily operated under quasi-steady-state assumptions and exhibited limited sensitivity to gradual or incipient leaks, where pressure deviations evolve slowly over time.

This work represents the most closely aligned prior research to the present study, demonstrating that leak-specific physics-informed graph learning can substantially enhance both detection reliability and spatial interpretability. Nevertheless, the framework predominantly emphasizes spatial correlations and does not explicitly model the temporal evolution of leak-induced dynamics. Moreover, the treatment of pressure transients and flow discontinuities remains approximate, particularly under evolving leak conditions and non-stationary operational regimes. These limitations motivate the development of a unified spatio-temporal graph-based physics-informed framework capable of jointly modeling network topology, temporal dynamics, and governing hydraulic laws for robust real-time leak detection and localization.

2.5. Benchmark Datasets and Evaluation Challenges

The availability of standardized datasets is crucial for reproducible research. **?** introduced LeakDB, a publicly available benchmark comprising EPANET-simulated scenarios for multiple WDN topologies (including Hanoi, Net1) with annotated leak events. The dataset facilitates comparative evaluation but

has been critiqued for relying solely on simulated data, potentially limiting model robustness under real-world operational uncertainties.

Comparative studies on BattLeDIM **?** and other benchmarks reveal persistent challenges: high false positive rates for incipient leaks, computational overhead in large-scale networks, and difficulty in achieving physically interpretable predictions under domain shifts. These gaps motivate the integration of physics-informed constraints within scalable GNN architectures.

2.6. Research Gaps and Motivation

Despite substantial progress, existing methodologies exhibit critical limitations:

- **Lack of Physical Interpretability:** Purely data-driven models **??** achieve high accuracies but fail to guarantee adherence to conservation laws, limiting generalization and trustworthiness.
- **Computational Scalability:** Hybrid approaches **?** and traditional PINNs **?** often require extensive hydraulic simulations or suffer from high computational costs in large networks.
- **Incipient Leak Detection:** Methods leveraging statistical anomaly detection **??** struggle with gradual leak growth due to subtle signal variations.

Addressing these gaps, the present study proposes a **Physics-Aware Artificial Neural Network (PA-ANN)** framework that embeds conservation of mass, the Hazen-Williams relation directly into a Feed-Forward Neural Network architecture. By using the LeakDB Hanoi dataset, we aim to achieve physically consistent computationally efficient leak detection, advancing beyond existing data-driven and hybrid methodologies.

3. What is a Physics-Aware Neural Network?

3.1. Physics-Informed Neural Networks (PINNs)

PINNs are neural networks that solve ***continuous partial differential equations (PDEs)*** that govern physical systems. This is done by embedding differential equations into the model training. The derivatives are calculated using ***automatic differentiation.*** The loss is defined as the PDE residual, evaluated at points in the continuous domain.

3.2. Physics-Aware Neural Networks (PANNs)

In comparison, as our data is discrete (at 30 minutes time-step), we use PANNs, here the model predicts the physical quantity along with its goal, and the predicted values are used to calculate the residuals of mass and energy using some formulae that seem appropriate for the data. The residuals (normalized) are added to the data loss for backprop, and the predicted values play a role in the classification part of the model. **Compared to PINNs, PANNs take significantly less training time and compute resources for a significant increase in performance over baseline models.**

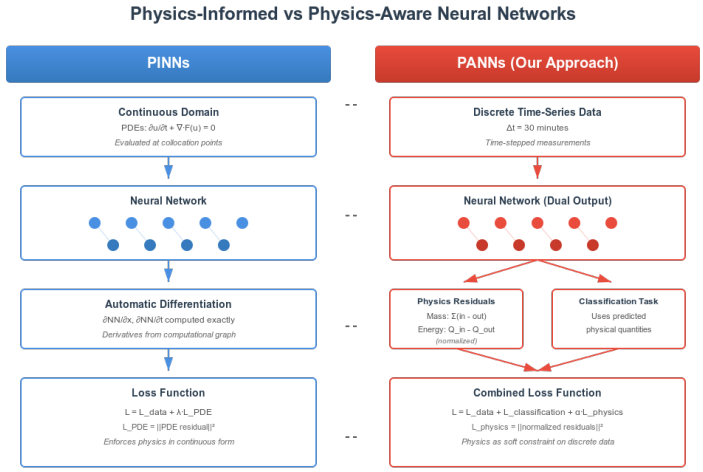

Figure 1: PINN vs PANN

3.3. Localization methods

3.4. Research Gap

4. Problem Formulation and System Overview

4.1. Problem Definition

4.2. Mathematical Formulation

4.3. System Overview Diagram

5. Dataset

5.1. Dataset Description

This study has used the LeakDB Hanoi dataset, which is made using the *Water Network Resilience Tool (WNTR)*. The dataset consists of a total of 999 scenarios, where each scenario is a different simulation of the Hanoi WDN for a year. The data points in each scenario are thirty minutes apart, which means there are 17,520 data points in one scenario, and a total of 17,502,480 data points in the entire dataset.

The Hanoi WDN consists of 32 nodes and 34 pipes connecting them, and the LeakDB dataset assumes pressure, demand and flow-rate sensors at these 32 nodes. Thus, each data point has 32 demand readings, 32 pipe readings and 34 flow-rate readings, making a total of 98 values per time-step.

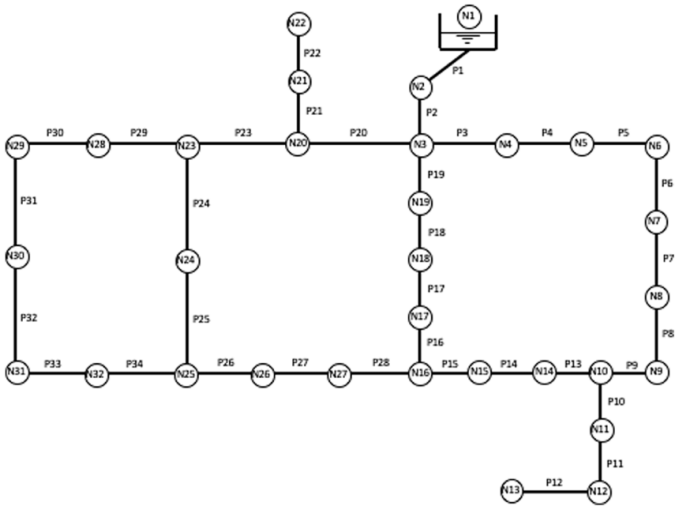

Figure 2: Figure of Hanoi WDN

The Data can be downloaded using the `water_benchmark_hub` Python module. There are two types of reading available, namely *raw_data* and *scada_data*. In this work, *raw_data* is used, as it also includes readings of demand.

Initially, each scenario is stored as a directory containing subfolders for *demand*, *pressure*, *flow*, and *leaks*, where leak labels are multi-class values ranging from 0 to 32. In addition, each scenario includes an *INP* file, which specifies the simulated length, diameter, and roughness coefficient of each pipe. For the learning task, the dataset is reformatted into a single `.csv` file per scenario, where leak labels are converted to a binary representation (leak / non-leak).

Leak localization is not learned directly by the model and is instead performed as a post-processing step using a rule-based system.

5.2. Exploratory Analysis of Data

Below is the heatmap of all 999 scenarios with 5% rows from each scenario selected randomly. The demand and pressure readings show a constant negative correlations, reflecting drop in pressure when demand increases or vice-versa. Flow readings show mixed correlations with demand and pressure, indicating the directional and topological effects of the network. These correlations support the use of dense models and also physics-aware models.

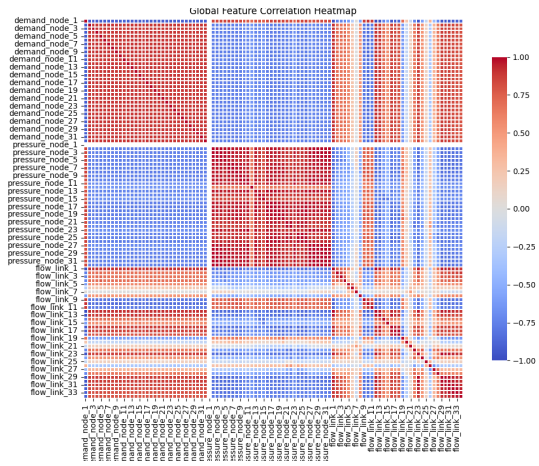

Figure 3: Correlation heatmap of Hanoi WDN

In the dataset, 43,11,716 rows have a leak state, while 1,31,90,764 rows a non-leak state, this means that that the ratio of non-leak data-points to leak points is **3.059**, making the dataset very imbalanced and a challenge for any neural network.

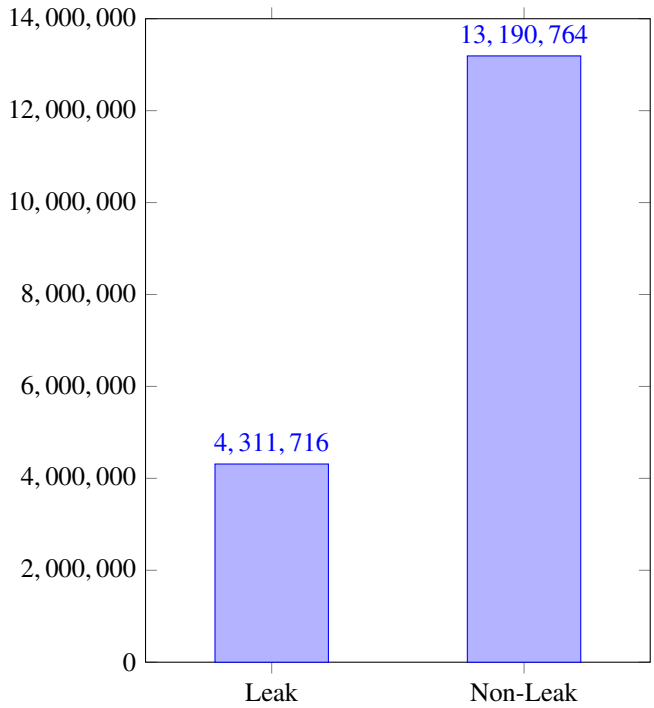

Figure 4: Distribution of leak and non-leak rows in the dataset.

A Scenario which has at least one leak state can be called a *leak scenario*, while a scenario which has no leak state in any timestamp can be called a *non_leak scenario*. Out of the 999 scenarios, 758 are *leak scenarios*, while 241 are *non_leak scenarios*. This ratio is also maintained while selecting the scenarios for training, validation, and testing.

Below is the distribution of leak instances across nodes from 0-31 (32 nodes).

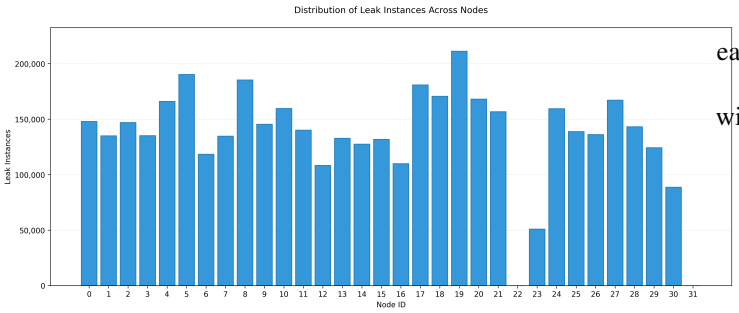

Figure 5: Caption

As we can see, node 19 has the highest leak instances, while node 22 and node 31 have no leaks. Using this, the rule-based system can be programmed to be biased more towards some nodes than others, in order for more correct node localizations.

5.3. Feature Engineering and Pre-Processing

In terms of feature engineering, two features have been added to each scenario for an increase in model performance. As the data does not have any missing or *NaN* values, filling in missing data is not needed and only data normalization has been performed in terms of pre-processing.

5.3.1. Adding columns

In the day time, due to usage of water both commercial and industrial, the general trend of pressure and flow is downward, while demand is higher and vice-versa during the night time. In order to prevent the models from reading the downward trend of pressure and flow and keep returning leak predictions, two columns, *sin_hour* and *cos_hour*, are added to each scenario, in order to add weight to the particular hour of the day when a leak occurs.

Let i denote row index of a scenario, Thus *step of day* $\boldsymbol{s}$ is defined as:

$$s_i = i \bmod 48$$

The temporal features are described as:

$$\text{sin_hour}_i = \sin\left(\frac{2\pi s_i}{48}\right)$$

$$\text{cos_hour}_i = \cos\left(\frac{2\pi s_i}{48}\right)$$

Because of these columns, a time step of 12 A.M. (sin_hour = 0 and cos_hour = 1) is given more bias by the model for a positive classification than 9 A.M. (sin_hour = 0.707 and cos_hour = -0.707) for similar readings of pressure, demand and flow.

5.3.2. Normalization

All sets (train, validation, and test) are normalized using the mean and standard deviation of the training set, such that the mean of each set is 0 and the standard deviation is 1, after normalization.

Let $\mathcal{D}_{\text{train}} = \{\mathbf{x}_1, \mathbf{x}_2, \ldots, \mathbf{x}_N\}$ denote the training dataset, where each feature vector $\mathbf{x}_i \in \mathbb{R}^F$.

The global mean and standard deviation are computed feature-wise over the training set as:

$$\boldsymbol{\mu} = \frac{1}{N}\sum_{i=1}^{N} \mathbf{x}_i, \tag{1}$$

$$\boldsymbol{\sigma} = \sqrt{\frac{1}{N}\sum_{i=1}^{N} \mathbf{x}_i^{\circ 2} - \boldsymbol{\mu}^{\circ 2}}, \tag{2}$$

where $\circ$ denotes element-wise operations.

To ensure numerical stability, the variance is lower bounded by a small constant ϵ:

$$\boldsymbol{\sigma} = \sqrt{\max\left(\frac{1}{N}\sum_{i=1}^{N} \mathbf{x}_i^{\circ 2} - \boldsymbol{\mu}^{\circ 2}, \epsilon\right)}. \tag{3}$$

Using the statistics computed from the training set, each feature is normalized as:

$$\tilde{\mathbf{x}} = \frac{\mathbf{x} - \boldsymbol{\mu}}{\boldsymbol{\sigma}}, \tag{4}$$

where $\tilde{\mathbf{x}}$ denotes the normalized feature vector.

The same training-set mean $\boldsymbol{\mu}$ and standard deviation $\boldsymbol{\sigma}$ are applied consistently to the training, validation, and test sets to avoid data leakage.

5.4. Windowing function

To give the model temporal context and features without the use of recurrent architectures such as LSTMs or GRUs, a sliding window based dataset construction is used. Temporal windows are extracted from each scenario using a fixed window size of 12, or 6 hours as each time-step is of 30 minutes. For each window, the target is the leak state (0/1) at the final time stamp of the window, allowing the model to learn temporal patterns preceding a leak event. As the target label is the leak state at the final time stamp, which is also the current time stamp, this means that the model **nowcasts** leakage prediction.

5.4.1. ANN

For the ANN models, each window is flattened into a one-dimensional vector as input to the network. A window of 12 time stamps with 100 features per time stamp is reshaped into a vector of dimension one and length 12×100. The mathematical representation is:

Let a scenario be represented as a multivariate time series

$$\mathbf{X} = \{\mathbf{x}_1, \mathbf{x}_2, \ldots, \mathbf{x}_T\}, \quad \mathbf{x}_t \in \mathbb{R}^F, \tag{5}$$

where F denotes the number of features at each time step.

Using a sliding window of length W, the k-th temporal window is defined as

$$\mathbf{X}^{(k)} = [\mathbf{x}_k, \mathbf{x}_{k+1}, \ldots, \mathbf{x}_{k+W-1}] \in \mathbb{R}^{W \times F}. \tag{6}$$

For the ANN model, each window is reshaped into a one-dimensional feature vector via vectorization:

$$\tilde{\mathbf{x}}^{(k)} = \mathrm{vec}\left(\mathbf{X}^{(k)}\right) \in \mathbb{R}^{W \cdot F}. \tag{7}$$

This flattened representation serves as the input to the fully connected network.

6. Physics Loss

All physical systems are governed by the conservation of mass and conservation of energy, and the Water Distribution Network is also assumed to be governed by these laws. The network is modeled as a directed graph, where vertices represent nodes, and edges represent pipes connecting nodes. *Mass conservation* is enforced at each node, and *Energy conservation* is enforced along each pipe. The equations used for these are discussed below.

6.1. Equations

For *Mass conservation*, the **Continuity Equation** is used. This means that the inflow at each node is equal to the sum of the demand at that node and the outflow at each node.

$$r_{\text{mass}}^{(i)} = \sum_{j \in \mathcal{E}_{\text{in}}(i)} Q_j - \sum_{j \in \mathcal{E}_{\text{out}}(i)} Q_j - D_i,$$

where Q_j denotes the flow through pipe j, D_i is the demand, and $\mathcal{E}_{\text{in}}(i)$ and $\mathcal{E}_{\text{out}}(i)$ represent incoming and outgoing pipes, respectively.

For *Energy conservation*, the *Hazen-Williams Equation* is used. The **Hazen-Williams Equation** defines the loss due to friction in a pipe; this frictional loss, along with the elevation head differences between the two ends of a pipe, gives the pressure drop. Since in the LeakDB dataset, in all scenarios, all nodes are kept at the same elevation, the pressure drop due to elevation differences is assumed to be negligible. The Hazen-Williams head loss for pipe k is given by

$$h_k = 10.67 \frac{L_k |Q_k|^{1.852}}{C_k^{1.852} D_k^{4.87}},$$

where L_k is the pipe length, D_k is the diameter, C_k is the roughness coefficient, and Q_k is the predicted flow.

6.2. Physics Loss Algorithm

The physics loss due to mass conservation is defined as:

$$\mathcal{L}_{\text{mass}} = \frac{1}{N} \sum_{i=1}^{N} \left(\frac{r_{\text{mass}}^{(i)}}{\bar{D}} \right)^2,$$

where $\bar{D}$ is the mean demand used for normalization.

The predicted pressure drop across a pipe is given as:

$$\Delta p_k = \mathrm{sign}(Q_k) \left(p_{\text{src}(k)} - p_{\text{dst}(k)} \right),$$

where $p_{\text{src}(k)}$ and $p_{\text{dst}(k)}$ are the predicted pressures at the source and destination nodes.

Thus, the loss due to energy conservation is defined as:

$$r_{\text{energy}}^{(k)} = \Delta p_k - h_k,$$

$$\mathcal{L}_{\text{energy}} = \frac{1}{M} \sum_{k=1}^{M} \left(\frac{r_{\text{energy}}^{(k)}}{\bar{p}} \right)^2,$$

where $\bar{p}$ is the mean predicted pressure used for normalization.

The total physics loss is a weighted average of the loss due to mass conservation and energy conservation, given as:

$$\mathcal{L}_{\text{phys}} = \lambda_{\text{mass}} \mathcal{L}_{\text{mass}} + \lambda_{\text{energy}} \mathcal{L}_{\text{energy}},$$

where λ_{mass} and λ_{energy} are both kept at 0.5.

6.3. Physics Loss Algorithm

The physics loss due to mass conservation is defined as:

$$\mathcal{L}_{\text{mass}} = \frac{1}{N} \sum_{i=1}^{N} \left(\frac{r_{\text{mass}}^{(i)}}{\bar{D}} \right)^2,$$

where $\bar{D}$ denotes a scale factor (e.g., mean absolute demand) for numerical stability.

The predicted pressure drop across a pipe is given as:

$$\Delta p_k = \mathrm{sign}(Q_k) \left(p_{\text{src}(k)} - p_{\text{dst}(k)} \right),$$

where $p_{\text{src}(k)}$ and $p_{\text{dst}(k)}$ are the predicted pressures at the source and destination nodes.

Thus, the loss due to energy conservation is defined as:

$$r_{\text{energy}}^{(k)} = \Delta p_k - h_k,$$

$$\mathcal{L}_{\text{energy}} = \frac{1}{M} \sum_{k=1}^{M} \left(\frac{r_{\text{energy}}^{(k)}}{\bar{p}} \right)^2,$$

where $\bar{p}$ denotes a scale factor (e.g., mean absolute predicted pressure) for numerical stability.

The total physics loss is a weighted average of the loss due to mass conservation and energy conservation, given as:

$$\mathcal{L}_{\text{phys}} = \lambda_{\text{mass}} \mathcal{L}_{\text{mass}} + \lambda_{\text{energy}} \mathcal{L}_{\text{energy}},$$

where λ_{mass} and λ_{energy} are hyperparameters; in our implementation we use $\lambda_{\text{mass}} = 1.0$ and $\lambda_{\text{energy}} = 0.1$ to emphasize mass conservation consistency.

6.4. Pseudocode and Training Integration

Algorithm 6.4 computes the physics loss given model predictions and pipe parameters. Algorithm 6.4 shows how it is integrated into the Physics-Aware ANN training loop. Predictions are denormalized before the physics loss is evaluated, since the conservation laws apply in physical units.

Algorithm 1: Physics Loss

Input: $\hat{p}$ (pred. pressure), $\hat{Q}$ (pred. flow), D (demand), L, D, C (pipe params), edge_index

Output: $\mathcal{L}_{\text{phys}}, \mathcal{L}_{\text{total}}$

1. Aggregate flow at nodes: $\text{flow_in}[i] \leftarrow \sum_{j:\text{dst}(j)=i} \hat{Q}_j$, $\text{flow_out}[i] \leftarrow \sum_{j:\text{src}(j)=i} \hat{Q}_j$
2. Mass residual: $r_{\text{mass}}^{(i)} \leftarrow \text{flow_in}[i] - \text{flow_out}[i] - D_i$
3. $\mathcal{L}_{\text{mass}} \leftarrow \text{mean}((r_{\text{mass}}/\bar{D})^2)$
4. Hazen-Williams head loss: $h_k \leftarrow 10.67 \cdot L_k |\hat{Q}_k|^{1.852} / (C_k^{1.852} D_k^{4.87})$
5. Predicted pressure drop: $\Delta p_k \leftarrow \text{sign}(\hat{Q}_k) \cdot (\hat{p}_{\text{src}(k)} - \hat{p}_{\text{dst}(k)})$
6. Energy residual: $r_{\text{energy}}^{(k)} \leftarrow \Delta p_k - h_k$
7. $\mathcal{L}_{\text{energy}} \leftarrow \text{mean}((r_{\text{energy}}/\bar{p})^2)$
8. $\mathcal{L}_{\text{phys}} \leftarrow \lambda_{\text{mass}} \mathcal{L}_{\text{mass}} + \lambda_{\text{energy}} \mathcal{L}_{\text{energy}}$
9. $\mathcal{L}_{\text{total}} \leftarrow \mathcal{L}_{\text{WBCE}} + \lambda_{\text{phys}} \mathcal{L}_{\text{phys}}$

Algorithm 2: Training Loop with Physics Loss

1. Set λ_{phys} by epoch: 0.01 if $e < 10$, 0.03 if $10 \le e < 20$, else 0.05
2. For each batch $(x_b, y_b, \text{scenario_ids}, D_b)$:
   (a) Load pipe params (L, D, C) for each scenario in batch
   (b) Denormalize D_b to physical units
   (c) Forward: $(\hat{y}, \hat{p}, \hat{Q}) \leftarrow \text{model}(x_b)$
   (d) Denormalize $\hat{p}$, $\hat{Q}$ to physical units
   (e) $\mathcal{L}_{\text{WBCE}} \leftarrow \text{weighted_BCE}(\hat{y}, y_b)$
   (f) $\mathcal{L}_{\text{total}} \leftarrow \text{PhysicsLoss}(\mathcal{L}_{\text{WBCE}}, \hat{p}, \hat{Q}, D_b, L, D, C, \text{edge_index}, \lambda_{\text{phys}})$
   (g) Backpropagate $\mathcal{L}_{\text{total}}$, update parameters

6.5. Integration into Model

To use this loss in the ANN model, the *.inp* files of each scenario are formatted into a *.csv* file containing the pipe length, diameter, and roughness coefficient of each pipe in the scenario, and this information is also integrated into the dataset class. As the node-pipe connection is constant for all scenarios, a constant *edge_list* is made and used. While training, the physics loss is added with the classification loss, and the result is used for backpropagation.

6.6. The Thought

Our idea with this loss was to constrain the model's latent space for such predictions that were physically possible and did not violate any of the equations, and that gradually the residual due to the difference between the true pressure and flow, and the predicted pressure and flow by the physics-aware models kept going closer and closer to 0. The predicted values of pressure and flow are also used in classification, discussed in more detail in the model architecture.

7. Model Architecture and Training Details

7.1. Architecture

7.1.1. Base ANN Model

The model here is a fully connected feed-forward network designed for classification of presence of leak on a single timestamp. The input layer consists of 1200 features, due to the Windowing function. It is passed through three hidden layers of 512, 256 and 64 neurons. The first two hidden layers employ Batch Normalization and ReLU activation, and a Dropout Layer (rate = 0.5) to prevent overfitting. The third hidden layer uses only ReLU activation. The output layer consists of a single neuron with a sigmoid activation function, for binary leak classification.

ANN Loss:

As discussed above, because the model is a binary classification model, and the class imbalance is severe, **Class-weighted binary cross-entropy** loss is used, where the loss is reweighted inversely proportional to class frequency.

Let N_0 denote the number of non-leak samples, and let N_1 denote the number of leak samples.

$$w_0 = \frac{N_0 + N_1}{2N_0}, \qquad w_1 = \frac{N_0 + N_1}{2N_1}$$

Let $y_i \in \{0, 1\}$ denote the ground-truth leak label for the i-th sample, and let $\hat{y}_i \in (0, 1)$ denote the corresponding predicted leak probability.

$$\mathcal{L}_{\text{WBCE}} = -\frac{1}{N} \sum_{i=1}^{N} \left[w_1 \, y_i \log(\hat{y}_i) + w_0 \, (1 - y_i) \log(1 - \hat{y}_i) \right]$$

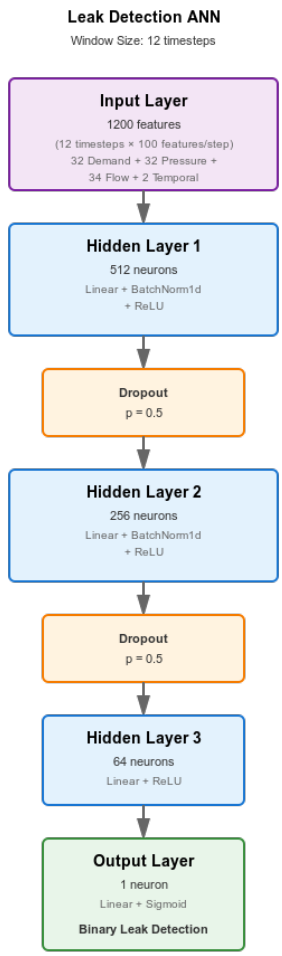

Figure 6: Architecture Diagram of ANN

7.1.2. Physics-Aware ANN Model

The model here is a fully connected feed-forward network designed for prediction of pressure and flow, and classficition of

presence of leak on a single timestamp, using those predictions. The input layer is 1200 features here too. It is passed through three hidden layers of 512, 256 and 64 neurons. All three hiden layers employ Batch Normalization, ReLU Activation and a Dropout Layer (rate = 0.5). The third hidden layer is connected to 2 layers, one with 32 neurons, and one with 34 neurons, to predict pressure and flow, respectively. The Leak head is consists of 2 layers, with the input layer being the 64 neurons of the third hidden layer, and 5 features that are derived from the predicted pressure and flow (mean of pressure and flow, standard deviation of pressure and flow, and the range of pressure), thus 69 features in total, which are then passed through a hidden layer of 64 neurons, where ReLU Activation and Dropout Layer is employed, and the output layer is a single neuron with a sigmoid activation function, for binary leak classification.

Physics-Aware ANN Loss:

The **Class-weighted binary cross-entropy**, also the **Classification loss** of the model is given by:

Let N_0 denote the number of non-leak samples, and let N_1 denote the number of leak samples.

$$w_0 = \frac{N_0 + N_1}{2N_0}, \qquad w_1 = \frac{N_0 + N_1}{2N_1}$$

Let $y_i \in \{0, 1\}$ denote the ground-truth leak label for the i-th sample, and let $\hat{y}_i \in (0, 1)$ denote the corresponding predicted leak probability.

$$\mathcal{L}_{\text{WBCE}} = -\frac{1}{N} \sum_{i=1}^{N} \left[w_1 \, y_i \log(\hat{y}_i) + w_0 \, (1 - y_i) \log(1 - \hat{y}_i) \right]$$

Let the physics loss, as discussed above in detail above, be represented as: $\mathcal{L}_{\text{phys}}$

Thus, the total loss of the model is given by:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{WBCE}} + \lambda_{\text{phys}} \, \mathcal{L}_{\text{phys}}.$$

The λ_{phys} follows this schedule:

$$\lambda_{\text{phys}}(e) = \begin{cases} 0.1, & e < 10, \\ 0.3, & 10 \le e < 20, \\ 0.5, & 20 \le e < 30, \end{cases}$$

This is done to make sure that the physics loss does not dominate the classification loss at the start.

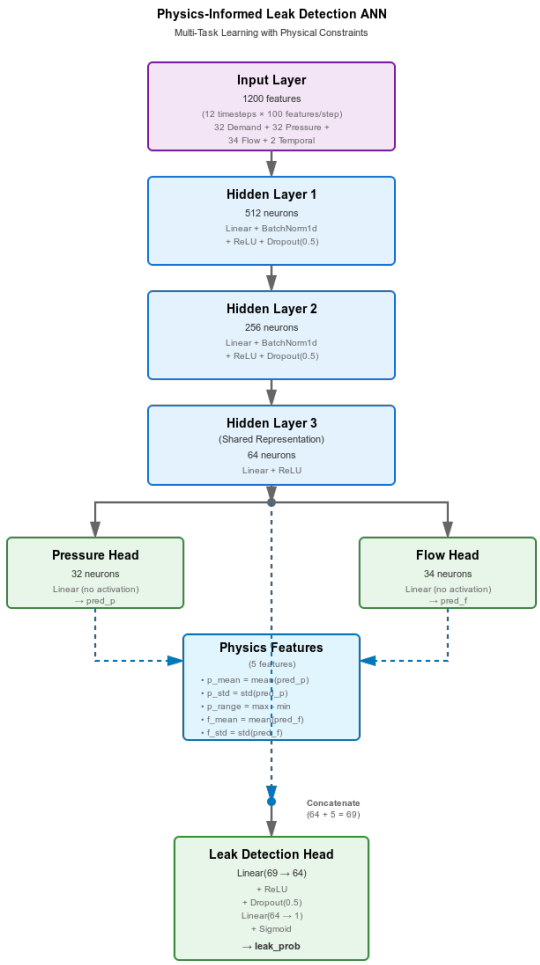

Figure 7: Architecture Diagram of Physics-Aware ANN

7.2. *Training strategy and Hyper-parameters*

8. Localization

8.1. *Introduction*

While the Deep Learning (ANN) module is responsible for the temporal detection of anomalies (identifying *when* a leak occurs), the spatial localization (identifying *where* the leak is located) is handled by a separate, physics-informed statistical engine. This hybrid approach leverages the strengths of both domains: the ANN captures complex, non-linear patterns to detect incipient events, while the localization module relies on hydraulic first principles—specifically, the analysis of pressure residuals relative to a dynamic baseline.

The localization logic operates on the hypothesis that a leak event introduces a localized pressure drop that is statistically significant compared to the node's historical variance. To ensure robustness against sensor noise and hydraulic transients, the system employs a "triggered snapshot" mechanism. Localization is not a continuous process; it is a conditional state activated only when the ANN confidence score exceeds a defined threshold

($\tau_{conf} > 0.75$). This ensures that the localization algorithm only processes frames where a high probability of a physical anomaly exists, thereby reducing false positives.

8.2. Algorithm

The localization algorithm proceeds in three distinct phases: Baseline Acquisition, Residual Standardization, and Temporal Aggregation.

8.2.1. 1. Dynamic Baseline Acquisition

During normal operation (when the ANN probability $P(leak) < 0.2$), the system maintains a rolling buffer of pressure readings for all N sensor nodes. Let $p_i(t)$ denote the pressure at node i at time t. The system continuously updates the mean (μ_i) and standard deviation (σ_i) for each sensor based on a window of stable steps W_{stable}. Upon the detection of a leak event at time $t_{trigger}$, the updating of this buffer is immediately halted. This freezes the "healthy" state of the network, providing a static reference frame against which the unfolding leak event is compared.

8.2.2. 2. Residual Standardization (Z-Score Analysis)

Once the system enters the active leak state, it calculates the raw pressure drop (residual) for every sensor at each time step t:

$$\Delta p_i(t) = \mu_i - p_i(t) \tag{8}$$

To account for sensor heterogeneity—where some nodes naturally exhibit higher variance than others—the raw drop is standardized into a Z-score. A minimum noise floor ($\epsilon \approx 0.1$) is applied to σ_i to prevent hypersensitivity in stable sensors:

$$Z_i(t) = \frac{\Delta p_i(t)}{\max(\sigma_i, \epsilon)} \tag{9}$$

A physics-based gate is applied to filter out negligible fluctuations. A sensor i is considered a candidate only if its drop exceeds a minimum hydraulic significance threshold ($\delta_{min} = 0.2$) and its statistical significance exceeds 3σ:

$$S_i(t) = \begin{cases} Z_i(t) & \text{if } \Delta p_i(t) > \delta_{min} \text{ and } Z_i(t) > 3.0 \\ 0 & \text{otherwise} \end{cases} \tag{10}$$

8.2.3. 3. Cumulative Evidence and Localization

To mitigate the impact of instantaneous sensor noise, the algorithm does not rely on a single time step. Instead, it employs a cumulative voting mechanism over the duration of the detected event ($t_{trigger}$ to $t_{current}$). The cumulative score C_i for each node is updated recursively:

$$C_i = \sum_{t=t_{trigger}}^{t_{current}} S_i(t) \tag{11}$$

The predicted leak location $\hat{L}$ is identified as the node with the maximum accumulated evidence at the end of the event window:

$$\hat{L} = \underset{i \in \{1,\ldots,N\}}{\operatorname{argmax}}(C_i) \tag{12}$$

This accumulation ensures that the localization decision converges towards the node that exhibits the most consistent and statistically significant pressure drop over the duration of the anomaly.

Algorithm 1 Hybrid Leak Detection and Localization Strategy

Require: Real-time Pressure Data stream P_t, Trained ANN Model $\mathcal{M}$
Ensure: Detected Leak Event E, Localized Node L_{est}

```
Initialize:
B ← ∅                                   ▷ Rolling Baseline Buffer
S ← {0, ..., 0}                         ▷ Cumulative Score Vector for N nodes
LeakActive ← False
while stream is active at time t do
    x_t ← Preprocess(P_{t−w:t})         ▷ Sliding Window Input
    p_leak ← M(x_t)                     ▷ Get Probability from ANN
    if p_leak < 0.2 and not LeakActive then
        Update(B, P_t)                  ▷ Update healthy baseline stats
        μ, σ ← CalculateStats(B)
    else if p_leak > 0.75 then
        LeakActive ← True               ▷ Trigger Event
    end if
    if LeakActive then
        for each sensor node i ∈ {1, ..., N} do
            Δ_i ← μ_i − P_{i,t}         ▷ Calculate Pressure Drop
            Z_i ← Δ_i / σ_i             ▷ Standardize (Z-Score)
            // Physics Gate: Ignore negligible drops
            if Δ_i > δ_min and Z_i > 3.0 then
                S_i ← S_i + Z_i         ▷ Accumulate Evidence
            end if
        end for
        L_est ← argmax_i(S)             ▷ Identify leading suspect
    end if
    if p_leak < 0.50 for k steps then
        LeakActive ← False              ▷ Reset System
        return L_est
    end if
end while
```

8.3. Metrics

9. Results and Discussion

9.1. Metrics

We use four standard metrics: **Accuracy**, **Precision**, **Recall**, and **F1-Score** to evaluate the performance of our classification models.

Accuracy measures the overall ratio of correctly classified instances.

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN} \tag{13}$$

$$\tag{14}$$

Precision measures how many positive predictions are also true

positives.

$$\text{Precision} = \frac{TP}{TP + FP} \tag{15}$$

$$\tag{16}$$

Recall measures how many of the true positives are actually predicted as true.

$$\text{Recall} = \frac{TP}{TP + FN} \tag{17}$$

$$\tag{18}$$

F1-Score is a harmonic mean of the Precision and Recall, acting as a single number that tells us about the performance of any model.

$$\text{F1-score} = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} \tag{19}$$

9.2. *Base Model Performance*

The ANN model is trained on 490 scenarios, with the validation and test sets being 105 scenarios each. The ANN model is trained for a set 30 epochs, and the model with the highest validation F1 is used for testing. The model is tested at a set threshold of **0.5**.

Highest Validation F1 for the ANN model is **0.7483**.

9.2.1. *ANN*

Model Metrics:

Accuracy	Precision	Recall	F1
0.9213	0.9143	0.7889	0.8470

Confusion Matrix:

	Predicted No Leak	Predicted Leak
Actual No Leak	1293741	37515
Actual Leak	107085	400104

Plot of Val F1 vs Epochs:

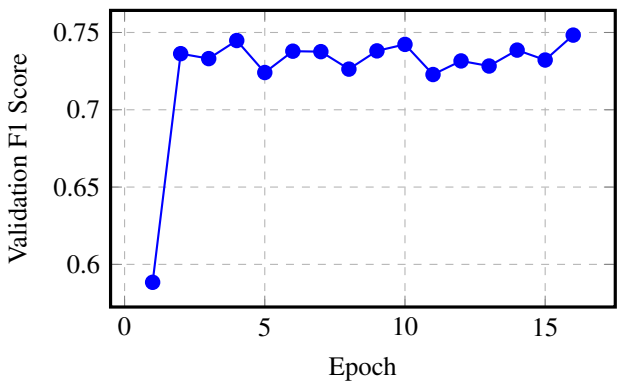

Figure 8: Validation F1 score across training epochs for ANN

9.3. *Physics-Aware Model Performance*

The ANN model is trained on 490 scenarios, with the validation and test sets being 105 scenarios each. The ANN model is trained for a set 30 epochs, and the model with the highest validation F1 is used for testing. The model is tested at a set threshold of **0.5**.

Highest Validation F1 for the ANN model is **0.7549**.

9.3.1. *ANN*

Model Metrics:

Accuracy	Precision	Recall	F1
0.9303	0.9659	0.7747	0.8598

Confusion Matrix:

	Predicted No Leak	Predicted Leak
Actual No Leak	1317369	13887
Actual Leak	114272	392917

Plot of Val F1 vs Epochs:

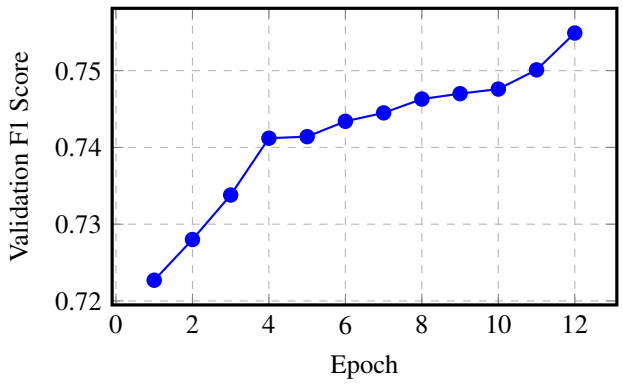

Figure 9: Validation F1 score across training epochs for Physics-Aware ANN

9.4. *Comparison among models and existing Benchmarks*

Comparison between ANN and Physics-Aware ANN:

Model	Accuracy	Precision	Recall	F1
ANN	0.9213	0.9143	**0.7889**	0.8470
Physics-Aware ANN	**0.9303**	**0.9659**	0.7747	**0.8598**

The Physics-Aware ANN gets an improvement of 0.98% in accuracy, 5.64% in precision, and 1.51% in F1 compared to the ANN, while having a decrease of 1.79% in recall. The Physics-Aware ANN reduces the number of false positives by 23,628 (62.98% reduction), at the cost of increasing the false negatives by 7,187 (6.71% increase). This indicates that the Physics-Aware ANN is more conservative while predicting leaks, issuing leak predictions only when they are physically consistent and more confident than the ANN, which has a more liberal behavior while predicting leaks.

In practical deployment, this trade-off means that the Physics-Aware ANN incurs a slight delay in detecting certain leak events, however, it also drastically lowers the number of false alarms, resulting in a significant reduction of un-necessary human intervention and monetary costs.

10. Why Classical Machine Learning Techniques Are Not Used

To assess whether classical machine learning methods could suffice for leak detection, we trained XGBoost classifiers on the same windowed dataset. Table 1 summarizes the results for N-estimators = 50 and 100, evaluated at threshold 0.5 and at the optimal threshold (ROC curve–derived, approximately 0.44). While XGBoost performs reasonably, it lags behind both the base ANN (F1 0.8470) and the Physics-Aware ANN (F1 0.8598). More importantly, XGBoost is purely data-driven and does not incorporate conservation of mass and energy. As a result, it cannot guarantee physically consistent predictions and tends toward higher false positive rates. The Physics-Aware ANN, by embedding hydraulic constraints into the loss, achieves better precision and lower false positives while maintaining strong recall, making it preferable for operational deployment.

Table 1: XGBoost on LeakDB Hanoi. N: trees; Thr.: threshold. ROC-AUC: 0.8705 (N=50), 0.8752 (N=100).

N	Thr.	Acc.	Prec.	Rec.	F1
50	0.50	0.8959	0.9519	0.6557	0.7765
50	0.44	0.8183	0.8695	0.7003	0.7758
100	0.50	0.9042	0.9567	0.6836	0.7974
100	0.44	0.8249	0.8905	0.7141	0.7926

11. Validity of Loss Function

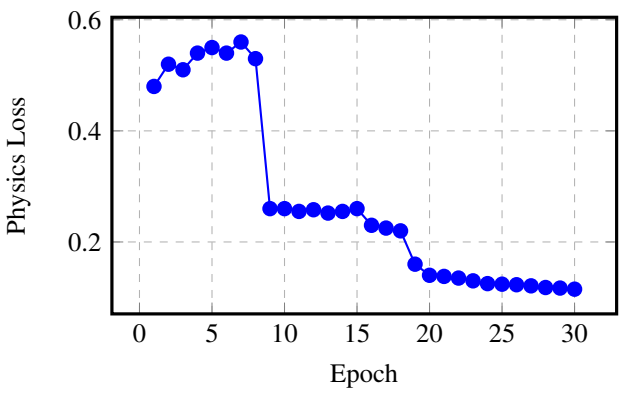

Figure 10: Physics loss convergence across training epochs

Any valid loss function used in optimization of a model should exhibit a decreasing trend during training if the model is able to learn the objective. Hence, similarly to how the weighted BCE Loss decreases across epochs as the model learns classification, the physics loss of the ANN model also decreases as it learns the underlying physics and the predicted pressure and flow values of the model leave smaller residuals across epochs.

12. Sensitivity Analysis

To evaluate model stability under distributional shifts, we performed a two-dimensional sensitivity analysis over demand perturbation $\alpha \in [-0.15, 0.15]$ and pressure–demand coupling strength $\beta \in [0.2, 1.0]$. Figure X and Figure Y present the F1-score surfaces for the Physics-informed ANN and the baseline ANN, respectively.

Both models achieve comparable performance under nominal conditions ($\alpha = 0$). However, as the magnitude of perturbation increases, distinct degradation patterns emerge. In particular, under negative demand shifts combined with stronger pressure sensitivity (i.e., higher β), the baseline ANN exhibits rapid performance deterioration, whereas the Physics-ANN demonstrates more gradual degradation.

Figure Z illustrates the performance gap $\Delta F_1 = F_{1,\text{Physics}} - F_{1,\text{ANN}}$. Positive regions indicate scenarios where physics-based constraints improve robustness. The advantage of the Physics-ANN becomes increasingly pronounced as β increases, suggesting that embedding physical relationships enhances stability under stronger pressure demand coupling.

Importantly, both models converge under extreme positive demand shifts, indicating that severe distribution shifts affect data-driven and physics-informed models in a similar manner. This observation highlights that physics constraints improve robustness within realistic operating regimes rather than acting as a universal correction mechanism.

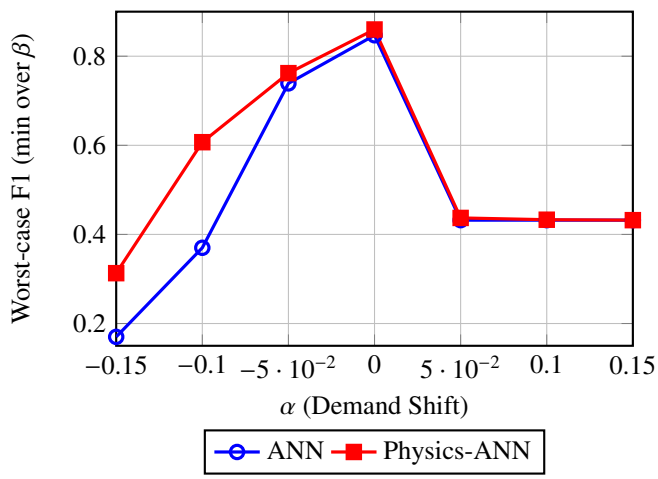

Figure 11: Worst-case robustness across demand perturbations. The Physics-informed ANN mitigates performance collapse under negative demand shifts.

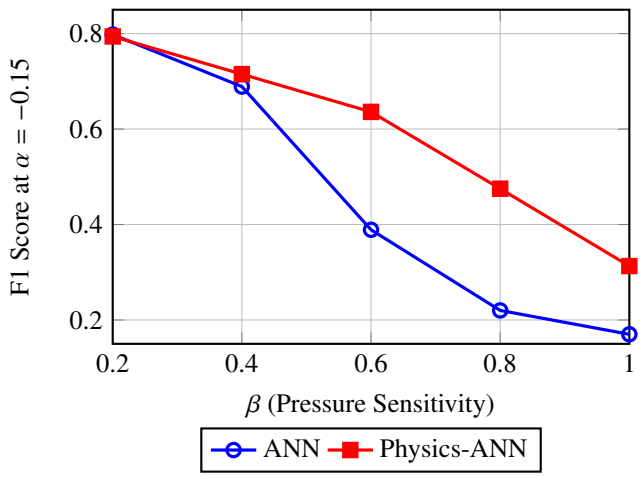

Figure 12: Performance under extreme demand reduction ($\alpha = -0.15$). The Physics-ANN exhibits significantly slower degradation as pressure sensitivity increases.

13. Localization Results

13.1. Quantitative Results

13.2. Qualitative Results

13.3. Comparison with Baselines

14. Limitations

Experiments were conducted on the LeakDB Hanoi benchmark, a widely used simulated dataset that allows controlled, reproducible evaluation. The Hanoi network (32 nodes, 34 pipes) serves as a standard testbed in the leak-detection literature. Training leveraged cloud compute for efficiency. The sensitivity analysis employs a tractable pressure–demand coupling to explore robustness under demand shifts; future work may incorporate full hydraulic simulation for finer-grained perturbation modeling.

15. Conclusion

We presented a Physics-Aware Artificial Neural Network (PAANN) for leak detection in Water Distribution Networks, embedding conservation of mass and the Hazen-Williams relation into the training objective. Experiments on the LeakDB Hanoi dataset show that the Physics-Aware ANN outperforms both a base ANN and XGBoost, achieving an F1 score of 0.8598 with a 62.98% reduction in false positives compared to the base ANN. By enforcing physical consistency, the model issues leak predictions only when they are hydraulically plausible, reducing false alarms and operational costs. We also introduced a sensitivity analysis framework to assess robustness under demand perturbations and discussed why classical machine learning methods such as XGBoost are less suitable for this task. Our results demonstrate that incorporating domain physics into neural network training improves detection reliability and supports the case for physics-aware approaches in critical infrastructure monitoring.

16. Future Work

Natural extensions include applying the Physics-Aware ANN to additional LeakDB topologies (e.g., Net1) and real-world SCADA datasets, incorporating EPANET-based perturbations in the sensitivity analysis, and exploring hybrid graph–physics architectures for larger networks. The rule-based localization component can be extended and benchmarked against alternative approaches. These directions build on the physics-aware framework presented here.

17. Acknowledgements

Each member of this group has contributed equally and with the same sense of enthusiasm and ownership for this project.
We would like to acknowledge Prof. Abhijeet Salunke for his mentorship, guidance and constant encouragement across the entire journey of the project.
We would also like to acknowledge Prof. Vasan Arunachalam from BITS Pilani, Hyderabad campus for solving a doubt of ours regarding the LeakDB dataset.

18. Appendix

18.1. Ablation Studies

We report ablations on three design choices: window size, temporal columns, and the physics-aware mechanism. All ablations use the base ANN unless otherwise noted; the physics-aware mechanism ablation compares two variants of the Physics-Aware ANN on a 100-scenario subset.

18.1.1. Window Size

Temporal context is provided by a sliding window. Table 2 and Figure 13 show that increasing window size from 1 to 12 improves F1 (0.617 → 0.680 → 0.703). A window of 12 (6 hours at 30-min resolution) captures daily patterns and yields the best performance.

Table 2: Ablation: window size.

Window	Acc.	Prec.	Rec.	F1
$W = 1$	0.8014	0.9195	0.4640	0.6168
$W = 6$	0.8271	0.9362	0.5344	0.6804
$W = 12$	0.8323	0.9009	0.5767	0.7033

18.1.2. Temporal Columns

The `sin_hour` and `cos_hour` features encode time-of-day. Table 3 and Figure 14 show that including them improves F1 from 0.686 to 0.703, indicating that the model benefits from explicit temporal conditioning.

Table 3: Ablation: temporal columns.

Setting	Acc.	Prec.	Rec.	F1
Without temporal	0.8258	0.9057	0.5518	0.6858
With temporal	0.8323	0.9009	0.5767	0.7033

18.1.3. Physics-Aware Mechanism

We ablate how physics is integrated into the Physics-Aware ANN. **Loss-only**: the physics loss is added to the classification loss, but the leak head receives only the shared hidden representation (no predicted pressure/flow). **Moderate coupling** (current): the leak head additionally receives 5 features derived from predicted pressure and flow (mean, std, range), as described in the model architecture. Table 4 and Figure 15 show that moderate coupling improves F1 from 0.70 to 0.72 on a 100-scenario subset, indicating that feeding physics-derived features into classification helps beyond the regularizing effect of the loss alone.

Table 4: Ablation: physics-aware mechanism (100 scenarios).

Variant	**Description**	**F1**
Loss only	Physics loss only; no pressure/flow in leak head	0.70
Moderate coupling	Predicted pressure/flow features in leak head	0.72